

0. Imports

```
[ ]: # read in data from tiingo
from pandas_datareader import data as pdr
# basic dataframe and array manipulation
import numpy as np
import pandas as pd
from datetime import datetime
# data visualization
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use("ggplot")
# keras sequential model pipeline
import tensorflow as tf
from tensorflow.keras.layers import Conv1D, MaxPooling1D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.layers import LSTM, ConvLSTM1D, TimeDistributed
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.regularizers import L1, L2
from tensorflow.keras.metrics import mean_absolute_error, mean_squared_error
from sklearn.metrics import r2_score
from tensorflow.keras.utils import plot_model
# for GridSearch CV
from keras.wrappers.scikit_learn import KerasRegressor
from sklearn.model_selection import GridSearchCV
# for timing training
import timeit
# plotting results: formating x axis date
import matplotlib.dates as mdates

[ ]: # get spy data from tiingo
with open('credentials/tiingo','r') as f:
    TIINGO_API_KEY = f.read().strip()
# get train data, 6 year
start=datetime(2010, 1, 1) # start day
end=datetime(2016, 12, 31) # end day
train_data = pdr.DataReader(["AAPL"], 'tiingo', start, end, api_key = TIINGO_API_KEY).
    loc["AAPL"]
# get test data, 3 year, skipping covid boom period (2020 spring)
start=datetime(2017, 1, 1) # start day
end=datetime(2019, 12, 31) # end day
val_data = pdr.DataReader(["AAPL"], 'tiingo', start, end, api_key = TIINGO_API_KEY).
    loc["AAPL"]
# get test data, 3 year, skipping covid boom period (2020 spring)
start=datetime(2020, 5, 1) # start day
end=datetime(2023, 5, 1) # end day
test_data = pdr.DataReader(["AAPL"], 'tiingo', start, end, api_key = TIINGO_API_KEY).
    loc["AAPL"]
```

1. Data Preprocessing

```
[ ]: window_size = 14
def window_data(df, window_size):
    X = []
    y = []
    for i in range(1, len(df) - window_size - 1, 1):
        first = df.iloc[i,2]
        temp = []
        temp2 = []
        for j in range(window_size):
            temp.append((df.iloc[i + j, 2] - first) / first)
            temp2.append((df.iloc[i + window_size, 2] - first) / first)
        X.append(np.array(temp).reshape(window_size, 1))
        y.append(np.array(temp2).reshape(1, 1))
    return X, y

X_train, y_train = window_data(train_data, window_size)
X_val, y_val = window_data(val_data, window_size)
X_test, y_test = window_data(test_data, window_size)

X_train = np.array(X_train)
X_val = np.array(X_val)
X_test = np.array(X_test)
y_train = np.array(y_train)
y_val = np.array(y_val)
y_test = np.array(y_test)

X_train = X_train.reshape(X_train.shape[0],1,window_size,1)
X_val = X_val.reshape(X_val.shape[0],1,window_size,1)
X_test = X_test.reshape(X_test.shape[0],1,window_size,1)
```

```
[ ]: print(X_train.shape)
print(X_val.shape)
print(X_test.shape)
```

```
(1746, 1, 14, 1)
(738, 1, 14, 1)
(739, 1, 14, 1)
```

2. Model construction

2.a Baseline model: FC-LSTM

2.a.(i) Hyperparameter Tuning

```
[ ]: def create_fc_lstm(dropout_rate=0.5, recurrent_dropout=0.5, lstm_units = 64):
    model = tf.keras.Sequential()
    model.add(LSTM(lstm_units, recurrent_dropout=recurrent_dropout,
        ↪return_sequences=True,
                    input_shape=(window_size, 1)))
    model.add(Dropout(dropout_rate))
```

```

    model.add(LSTM(lstm_units, recurrent_dropout=recurrent_dropout, activation='relu',
↳return_sequences=True))
    model.add(Dropout(dropout_rate))
    model.add(LSTM(lstm_units, recurrent_dropout=recurrent_dropout, activation='relu',
↳return_sequences=True))
    model.add(Dropout(dropout_rate))
    model.add(Flatten())
    model.add(Dense(1))
    model.compile(optimizer='adam', loss='mae', metrics=['mse', 'mae'])
    return model
# GridSearch CV
param_grid = {
    'dropout_rate': [0.3, 0.5],
    'recurrent_dropout': [0.3, 0.5],
    'lstm_units': [32, 64, 128],
}
# create KerasRegressor and GridSearchCV objects
early_stopping = EarlyStopping(monitor='val_loss', patience=5)
tf.keras.backend.clear_session()
model = KerasRegressor(build_fn=create_fc_lstm, epochs=100, batch_size=32,
    verbose=0, callbacks=[early_stopping])
grid = GridSearchCV(estimator=model, param_grid=param_grid, cv=3, verbose=0)

# perform Grid Search CV
grid_result = grid.fit(X_train.reshape((len(X_train), X_train.shape[2], X_train.
↳shape[3])), y_train,
    validation_data=(X_val.reshape((len(X_val), X_val.shape[2],
↳X_val.shape[3])), y_val))
fc_lstm_best_params = grid_result.best_params_

# print results
print(f"Best params: {grid_result.best_params_}")

```

/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:24: DeprecationWarning: KerasRegressor is deprecated, use Sci-Keras (<https://github.com/adriangb/scikeras>) instead. See <https://www.adriangb.com/scikeras/stable/migration.html> for help migrating.

Best params: {'dropout_rate': 0.3, 'lstm_units': 128, 'recurrent_dropout': 0.3}

2.a.(ii) Training the model

```

[ ]: tf.keras.backend.clear_session()
fc_lstm = create_fc_lstm(dropout_rate=fc_lstm_best_params['dropout_rate'],
    recurrent_dropout=fc_lstm_best_params['recurrent_dropout'],
    lstm_units = fc_lstm_best_params['lstm_units'])
fc_lstm._name = "FC-LSTM"
early_stopping = EarlyStopping(monitor='val_loss', patience=5)
fc_lstm.summary()

```

Model: "FC-LSTM"

Layer (type)	Output Shape	Param #
=====		

lstm (LSTM)	(None, 14, 128)	66560
dropout (Dropout)	(None, 14, 128)	0
lstm_1 (LSTM)	(None, 14, 128)	131584
dropout_1 (Dropout)	(None, 14, 128)	0
lstm_2 (LSTM)	(None, 14, 128)	131584
dropout_2 (Dropout)	(None, 14, 128)	0
flatten (Flatten)	(None, 1792)	0
dense (Dense)	(None, 1)	1793

```

=====
Total params: 331,521
Trainable params: 331,521
Non-trainable params: 0
-----

```

```


[ ]: start_time = timeit.default_timer()
fc_lstm_history = fc_lstm.fit(X_train.reshape((len(X_train), X_train.shape[2], X_train.
↳shape[3])),y_train,
                                validation_data=(X_val.reshape((len(X_val), X_val.
↳shape[2], X_val.shape[3])),y_val),
                                epochs=100,batch_size=32, verbose=0,↳
↳callbacks=[early_stopping])
fc_lstm_elapsed = timeit.default_timer() - start_time
plt.plot(fc_lstm_history.history['loss'], label='train loss')
plt.plot(fc_lstm_history.history['val_loss'], label='val loss')
plt.xlabel("epoch")
plt.ylabel("Loss")
plt.legend()

```

```

[ ]: <matplotlib.legend.Legend at 0x15e0713c8>

```



2.b CNN-LSTM

2.b.(i) Hyperparameter tuning

```
[ ]: # Use Grid Search to Tune drop out
def create_cnn_lstm(dropout_rate=0.5, recurrent_dropout=0.5, cnn_units=[32, 64, 128],
                    lstm_units = 64, kernel_size=5):
    model = tf.keras.Sequential()
    # CNN layers
    model.add(TimeDistributed(Conv1D(cnn_units[0], kernel_size=kernel_size,
    activation='relu', input_shape=(None,
    window_size, 1))))
    model.add(TimeDistributed(Conv1D(cnn_units[1], kernel_size=kernel_size,
    activation='relu')))
    model.add(TimeDistributed(Conv1D(cnn_units[2], kernel_size=kernel_size,
    activation='relu')))
    model.add(TimeDistributed(MaxPooling1D(2)))
    model.add(TimeDistributed(Flatten()))
    # cnn_lstm.add(Dense(5, kernel_regularizer=L2(0.01)))
    # LSTM layers
    model.add(LSTM(lstm_units, recurrent_dropout=recurrent_dropout,
    return_sequences=True))
    model.add(Dropout(dropout_rate))
    model.add(LSTM(lstm_units, recurrent_dropout=recurrent_dropout,
    return_sequences=True))
    model.add(Dropout(dropout_rate))
    model.add(LSTM(lstm_units, recurrent_dropout=recurrent_dropout,
    return_sequences=False))
    model.add(Dropout(dropout_rate))
```

```

# FC layers, output
model.add(Dense(1, activation='linear'))
model.compile(optimizer='adam', loss='mae', metrics=['mse', 'mae'])
return model
# GridSearch CV
param_grid = {
    'dropout_rate': [0.3, 0.5],
    'recurrent_dropout': [0.3, 0.5],
    'cnn_units': [[32, 64, 128], [64, 64, 128]],
    'lstm_units': [32, 64, 128],
    'kernel_size': [3, 5]
}
# create KerasRegressor and GridSearchCV objects
early_stopping = EarlyStopping(monitor='val_loss',patience=5)
tf.keras.backend.clear_session()
model = KerasRegressor(build_fn=create_cnn_lstm, epochs=100, batch_size=32,
                        verbose=0, callbacks=[early_stopping])
grid = GridSearchCV(estimator=model, param_grid=param_grid, cv=3, verbose=0)

# perform Grid Search CV
grid_result = grid.fit(X_train, y_train, validation_data=(X_val, y_val))
cnn_lstm_best_params = grid_result.best_params_

# print results
print(f"Best params: {grid_result.best_params_}")

```

<ipython-input-12-be7f2ccfb87d>:35: DeprecationWarning: KerasRegressor is deprecated, use Sci-Keras (<https://github.com/adriangb/scikeras>) instead. See <https://www.adriangb.com/scikeras/stable/migration.html> for help migrating.

```

model = KerasRegressor(build_fn=create_cnn_lstm, epochs=100, batch_size=32,

Best params: {'cnn_units': [32, 64, 128], 'dropout_rate': 0.3, 'kernel_size': 3,
'lstm_units': 128, 'recurrent_dropout': 0.5}

```

2.b.(ii) Training the model

```

[ ]: tf.keras.backend.clear_session()
cnn_lstm = create_cnn_lstm(dropout_rate=cnn_lstm_best_params['dropout_rate'],
                           recurrent_dropout=cnn_lstm_best_params['recurrent_dropout'],
                           cnn_units=cnn_lstm_best_params['cnn_units'],
                           lstm_units=cnn_lstm_best_params['lstm_units'],
                           kernel_size=cnn_lstm_best_params['kernel_size'])
cnn_lstm._name = "CNN-LSTM"
early_stopping = EarlyStopping(monitor='val_loss',patience=5)
cnn_lstm.build(input_shape=(len(X_train),None, window_size, 1))
cnn_lstm.summary()

```

Model: "CNN-LSTM"

Layer (type)	Output Shape	Param #
time_distributed (TimeDistrib	(1746, None, 12, 32)	128
ibuted)		

time_distributed_1 (TimeDistributed)	(1746, None, 10, 64)	6208
time_distributed_2 (TimeDistributed)	(1746, None, 8, 128)	24704
time_distributed_3 (TimeDistributed)	(1746, None, 4, 128)	0
time_distributed_4 (TimeDistributed)	(1746, None, 512)	0
lstm (LSTM)	(1746, None, 128)	328192
dropout (Dropout)	(1746, None, 128)	0
lstm_1 (LSTM)	(1746, None, 128)	131584
dropout_1 (Dropout)	(1746, None, 128)	0
lstm_2 (LSTM)	(1746, 128)	131584
dropout_2 (Dropout)	(1746, 128)	0
dense (Dense)	(1746, 1)	129

```

=====
Total params: 622,529
Trainable params: 622,529
Non-trainable params: 0
-----

```

```


[ ]: start_time = timeit.default_timer()
      cnn_lstm_history = cnn_lstm.fit(X_train, y_train, validation_data=(X_val,y_val),
                                     epochs=100,batch_size=32, verbose=0, callbacks=[early_stopping])
      cnn_lstm_elapsed = timeit.default_timer() - start_time
      plt.plot(cnn_lstm_history.history['loss'], label='train loss')
      plt.plot(cnn_lstm_history.history['val_loss'], label='val loss')
      plt.xlabel("epoch")
      plt.ylabel("Loss")
      plt.legend()

```

```

[ ]: <matplotlib.legend.Legend at 0x159512780>

```



2.c ConVLSTM

2.c.(i) Hyperparameter tuning

```
[ ]: def create_convlstm(dropout_rate=0.5, recurrent_dropout=0.5, units=[64,64,128],
    ↪kernel_size=5):
    model = tf.keras.Sequential()
    model.add(ConvLSTM1D(filters=units[0], recurrent_dropout=recurrent_dropout,
    ↪kernel_size=kernel_size,
        padding="same", input_shape=(None, window_size, 1),
    ↪return_sequences=True))
    model.add(Dropout(dropout_rate))
    model.add(ConvLSTM1D(filters=units[1], recurrent_dropout=recurrent_dropout,
    ↪kernel_size=kernel_size,
        padding="same", return_sequences=True))
    model.add(Dropout(dropout_rate))
    model.add(ConvLSTM1D(filters=units[2], recurrent_dropout=recurrent_dropout,
    ↪kernel_size=kernel_size,
        padding="same", return_sequences=False))
    model.add(Dropout(dropout_rate))
    model.add(Flatten())
    model.add(Dense(1))
    model.compile(optimizer='adam', loss='mae', metrics=['mse', 'mae'])
    return model
# GridSearch CV
param_grid = {
    'dropout_rate': [0.3, 0.5],
    'recurrent_dropout': [0.3, 0.5],
    'units': [[128, 64, 32], [32, 64, 128], [64, 64, 64]],
```



```

    'kernel_size': [3, 5]
}
# create KerasRegressor and GridSearchCV objects
early_stopping = EarlyStopping(monitor='val_loss',patience=5)
model = KerasRegressor(build_fn=create_conv_lstm, epochs=100, batch_size=32,
                        verbose=0, callbacks=[early_stopping])
grid = GridSearchCV(estimator=model, param_grid=param_grid, cv=3, verbose=0)

# perform Grid Search CV
grid_result = grid.fit(X_train, y_train, validation_data=(X_val, y_val))
conv_lstm_best_params = grid_result.best_params_

# print results
print(f"Best params: {grid_result.best_params_}")

```

<ipython-input-15-e965e00549af>:25: DeprecationWarning: KerasRegressor is deprecated, use Sci-Keras (<https://github.com/adriangb/scikeras>) instead. See <https://www.adriangb.com/scikeras/stable/migration.html> for help migrating.

```
model = KerasRegressor(build_fn=create_conv_lstm, epochs=100, batch_size=32,
```

```
Best params: {'dropout_rate': 0.3, 'kernel_size': 5, 'recurrent_dropout': 0.5, 'units':
[128, 64, 32]}
```

2.c.(ii) Training the model

```

[ ]: tf.keras.backend.clear_session()
conv_lstm = create_conv_lstm(dropout_rate=conv_lstm_best_params['dropout_rate'],
                             recurrent_dropout=conv_lstm_best_params['recurrent_dropout'],
                             units=conv_lstm_best_params['units'],
                             kernel_size=conv_lstm_best_params['kernel_size'])
conv_lstm._name = "ConVLSTM"
early_stopping = EarlyStopping(monitor='val_loss',patience=5)
conv_lstm.build(input_shape=(len(X_train),None, 30, 1))
conv_lstm.summary()

```

Model: "ConVLSTM"


Layer (type)	Output Shape	Param #
conv_lstm1d (ConvLSTM1D)	(None, None, 14, 128)	330752
dropout (Dropout)	(None, None, 14, 128)	0
conv_lstm1d_1 (ConvLSTM1D)	(None, None, 14, 64)	246016
dropout_1 (Dropout)	(None, None, 14, 64)	0
conv_lstm1d_2 (ConvLSTM1D)	(None, 14, 32)	61568
dropout_2 (Dropout)	(None, 14, 32)	0
flatten (Flatten)	(None, 448)	0

dense (Dense) (None, 1) 449

```
=====
Total params: 638,785
Trainable params: 638,785
Non-trainable params: 0
-----
```

```
[ ]: start_time = timeit.default_timer()
convlstm_history = convlstm.fit(X_train, y_train,
    validation_data=(X_val,y_val),epochs=100,
                                batch_size=32, verbose=0, callbacks=[early_stopping])
convlstm_elapsed = timeit.default_timer() - start_time
plt.plot(convlstm_history.history['loss'], label='train loss')
plt.plot(convlstm_history.history['val_loss'], label='val loss')
plt.xlabel("epoch")
plt.ylabel("Loss")
plt.legend()
```

```
[ ]: <matplotlib.legend.Legend at 0x161f34ac8>
```



3. Testing Phase

```
[ ]: def plot_stock_prediction(window,model,X_test):
    test_result = model.evaluate(X_test, y_test)
    predicted = model.predict(X_test)
    true_Y = test_data['low'].iloc[(window+1):-1]
    true_Y_prev = test_data['low'].iloc[1:-1-window]
    df_test_time = true_Y.index.date
```

```

pred_Y = (np.array(predicted[:,0])*true_Y_prev)+true_Y_prev
print("R2 score of ",model._name,": ",r2_score(y_test.reshape(-1,1), predicted))
# plot test data
%matplotlib inline
fig, ax = plt.subplots(figsize=(30, 8))
monthly_locator = mdates.MonthLocator()
half_year_locator = mdates.MonthLocator(interval=1)
year_month_formatter = mdates.DateFormatter("%Y-%m")
ax.xaxis.set_major_locator(half_year_locator)
ax.xaxis.set_minor_locator(monthly_locator)
ax.xaxis.set_major_formatter(year_month_formatter)
ax.plot(np.array(df_test_time), true_Y,color = 'red', label = 'Real Stock Price')
ax.plot(np.array(df_test_time), pred_Y,color = 'green', label = 'Predicted Stock_
↪Price')
ax.title.set_text(f'Stock Price Prediction, \
MAE = \$ {round(float(mean_absolute_error(pred_Y,true_Y)),8)}, \
MSE = \$ {round(float(mean_squared_error(true_Y,pred_Y),8)}')
ax.legend(loc='upper left')
fig.autofmt_xdate()

```

3.a FC-LSTM

```


[ ]: plot_stock_prediction(window_size,fc_lstm,X_test.reshape((len(X_test), X_test.
↪shape[2], X_test.shape[3])))

```

24/24 [=====] - 1s 40ms/step - loss: 0.0241 - mse: 0.0020 - mae: 0.0241

24/24 [=====] - 1s 34ms/step

R2 score of FC-LSTM : 0.8740441414805302



```

[ ]: print("Model training time:",fc_lstm_elapsed)

```

Model training time: 94.75706969799921

3.b CNN-LSTM

```


[ ]: plot_stock_prediction(window_size,cnn_lstm,X_test)

```

24/24 [=====] - 0s 11ms/step - loss: 0.0160 - mse: 0.0014 - mae: 0.0160

24/24 [=====] - 0s 10ms/step

R2 score of CNN-LSTM : 0.9139324437237274



```
[ ]: print("Model training time:",cnn_lstm_elapsed)
```

Model training time: 35.57562987100005


3.c ConVLSTM

```
[ ]: plot_stock_prediction(window_size,convlstm,X_test)
```

24/24 [=====] - 1s 31ms/step - loss: 0.0157 - mse: 0.0013 - mae: 0.0157

24/24 [=====] - 1s 31ms/step

R2 score of ConVLSTM : 0.9177376534570053



```
[ ]: print("Model training time:",convlstm_elapsed)
```

Model training time: 161.20439217799958